

car-prediction

November 5, 2023

Import Required Libraries

```
[50]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import Lasso
from sklearn import metrics
```

Data Collection and processing

```
[51]: from google.colab import files
```

```
uploaded = files.upload()
```

<IPython.core.display.HTML object>

Saving car.csv to car (1).csv

```
[52]: car_dataset = pd.read_csv('car.csv')
```

```
[53]: car_dataset.head()
```

```
[53]: Car_Name  Year  Selling_Price  Present_Price  Driven_kms  Fuel_Type  \
0    ritz    2014           3.35           5.59       27000    Petrol
1    sx4    2013           4.75           9.54       43000    Diesel
2    ciaz    2017           7.25           9.85        6900    Petrol
3  wagon r    2011           2.85           4.15        5200    Petrol
4   swift    2014           4.60           6.87       42450    Diesel
```

```
    Selling_type  Transmission  Owner
0      Dealer      Manual      0
1      Dealer      Manual      0
2      Dealer      Manual      0
3      Dealer      Manual      0
4      Dealer      Manual      0
```

```
[54]: car_dataset.shape
```

```
[54]: (301, 9)
```

```
[55]: car_dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 301 entries, 0 to 300
Data columns (total 9 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   Car_Name        301 non-null    object
 1   Year            301 non-null    int64
 2   Selling_Price   301 non-null    float64
 3   Present_Price   301 non-null    float64
 4   Driven_kms      301 non-null    int64
 5   Fuel_Type       301 non-null    object
 6   Selling_type    301 non-null    object
 7   Transmission    301 non-null    object
 8   Owner           301 non-null    int64
dtypes: float64(2), int64(3), object(4)
memory usage: 21.3+ KB
```

```
[56]: car_dataset.isnull().sum()
```

```
[56]: Car_Name      0
      Year         0
      Selling_Price 0
      Present_Price 0
      Driven_kms    0
      Fuel_Type     0
      Selling_type  0
      Transmission  0
      Owner         0
      dtype: int64
```

```
[57]: categorical_columns = ['Car_Name', 'Fuel_Type', 'Selling_type', 'Transmission',
                             ↪ 'Owner']

for col in categorical_columns:
    print(car_dataset[col].value_counts())
```

```
city                26
corolla altis       16
verna               14
fortuner            11
brio                10
..
Honda CB Trigger     1
Yamaha FZ S          1
```

```

Bajaj Pulsar 135 LS      1
Activa 4g                1
Bajaj Avenger Street 220 1
Name: Car_Name, Length: 98, dtype: int64
Petrol      239
Diesel      60
CNG         2
Name: Fuel_Type, dtype: int64
Dealer      195
Individual   106
Name: Selling_type, dtype: int64
Manual      261
Automatic    40
Name: Transmission, dtype: int64
0          290
1           10
3           1
Name: Owner, dtype: int64

```

One-Hot Encoding

```

[58]: categorical_columns = ['Fuel_Type', 'Selling_type', 'Transmission']

car_dataset_encoded = pd.get_dummies(car_dataset, columns=categorical_columns,
↳drop_first=True)

```

```

[59]: car_dataset_encoded

```

```

[59]:   Car_Name  Year  Selling_Price  Present_Price  Driven_kms  Owner  \
0    ritz    2014         3.35         5.59       27000      0
1    sx4    2013         4.75         9.54       43000      0
2    ciaz    2017         7.25         9.85        6900      0
3  wagon r    2011         2.85         4.15        5200      0
4    swift    2014         4.60         6.87       42450      0
..    ...    ...         ...         ...         ...         ...
296  city    2016         9.50        11.60       33988      0
297  brio    2015         4.00         5.90       60000      0
298  city    2009         3.35        11.00       87934      0
299  city    2017        11.50        12.50        9000      0
300  brio    2016         5.30         5.90        5464      0

      Fuel_Type_Diesel  Fuel_Type_Petrol  Selling_type_Individual  \
0                   0                   1                       0
1                   1                   0                       0
2                   0                   1                       0
3                   0                   1                       0
4                   1                   0                       0
..                  ...                 ...                     ...

```

296	1	0	0
297	0	1	0
298	0	1	0
299	1	0	0
300	0	1	0

Transmission_Manual	
0	1
1	1
2	1
3	1
4	1
..	...
296	1
297	1
298	1
299	1
300	1

[301 rows x 10 columns]

Label Encoding:

```
[24]: from sklearn.preprocessing import LabelEncoder

categorical_columns = ['Car_Name']

label_encoder = LabelEncoder()

car_dataset['Car_Name'] = label_encoder.fit_transform(car_dataset['Car_Name'])
```

```
[25]: car_dataset.head()
```

```
[25]:   Car_Name  Year  Selling_Price  Present_Price  Driven_kms  Fuel_Type  \
0      90    2014           3.35           5.59       27000    Petrol
1      93    2013           4.75           9.54       43000    Diesel
2      68    2017           7.25           9.85        6900    Petrol
3      96    2011           2.85           4.15        5200    Petrol
4      92    2014           4.60           6.87       42450    Diesel
```

	Selling_type	Transmission	Owner
0	Dealer	Manual	0
1	Dealer	Manual	0
2	Dealer	Manual	0
3	Dealer	Manual	0
4	Dealer	Manual	0

Splitting the data and Target

[60]:

[61]: `print(X)`

	Year	Present_Price	Driven_kms	Owner	Fuel_Type_Diesel	\
0	2014	5.59	27000	0	0	
1	2013	9.54	43000	0	1	
2	2017	9.85	6900	0	0	
3	2011	4.15	5200	0	0	
4	2014	6.87	42450	0	1	
..	
296	2016	11.60	33988	0	1	
297	2015	5.90	60000	0	0	
298	2009	11.00	87934	0	0	
299	2017	12.50	9000	0	1	
300	2016	5.90	5464	0	0	

	Fuel_Type_Petrol	Selling_type_Individual	Transmission_Manual
0	1	0	1
1	0	0	1
2	1	0	1
3	1	0	1
4	0	0	1
..
296	0	0	1
297	1	0	1
298	1	0	1
299	0	0	1
300	1	0	1

[301 rows x 8 columns]

[62]: `print(Y)`

0	3.35
1	4.75
2	7.25
3	2.85
4	4.60
...	
296	9.50
297	4.00
298	3.35
299	11.50
300	5.30

Name: Selling_Price, Length: 301, dtype: float64

Splitting Training and Test Data

Model Training

1. Linear Regression

[67]:

Model Evaluation

[77]:

```
from sklearn.model_selection import train_test_split

X = car_dataset_encoded.drop(['Car_Name', 'Selling_Price'], axis=1)
Y = car_dataset_encoded['Selling_Price']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1,
                                                    random_state=2)
```

[78]:

```
X_train
```

[78]:

	Year	Present_Price	Driven_kms	Owner	Fuel_Type_Diesel	\
204	2015	4.430	28282	0	0	
249	2016	7.600	17000	0	0	
277	2015	13.600	21780	0	0	
194	2008	0.787	50000	0	0	
244	2013	9.400	49000	0	1	
..	
75	2015	6.800	36000	0	0	
22	2011	8.010	50000	0	0	
72	2013	18.610	56001	0	0	
15	2016	10.790	43000	0	1	
168	2013	0.730	12000	0	0	
	Fuel_Type_Petrol	Selling_type_Individual	Transmission_Manual			
204	1	0	1			
249	1	0	1			
277	1	0	1			
194	1	1	1			
244	0	0	1			
..			
75	1	0	1			
22	1	0	0			
72	1	0	1			
15	0	0	1			
168	1	1	1			

[270 rows x 8 columns]

[80]:


```
lin_reg_model = LinearRegression()
lin_reg_model.fit(X_train, y_train)
```

```
training_data_prediction = lin_reg_model.predict(X_train)
```

```
[84]: error_score = metrics.r2_score(y_train, training_data_prediction)
      print("R squared Error :", error_score)
```

R squared Error : 0.8823856405331196

```
[85]: plt.scatter(y_train, training_data_prediction)
      plt.xlabel("Actual Price")
      plt.ylabel("Predicted Price")
      plt.title("Actual Proces Vs Predicted Proces")
      plt.show()
```




```
[86]: test_data_prediction = lin_reg_model.predict(X_test)
```

```
[87]: error_score = metrics.r2_score(y_test, test_data_prediction)
      print("R squared Error : ", error_score)
```

R squared Error : 0.8694567179819735

```
[88]: plt.scatter(y_test, test_data_prediction)
plt.xlabel('Actual Price')
plt.ylabel("Predicted Price")
plt.title("Actual Procees vs Predicted Prices")
plt.show()
```



```
[89]: lass_reg_model=Lasso()
```

```
[90]: lass_reg_model.fit(X_train, y_train)
```


```
[90]: Lasso()
```

```
[97]: training_data_prediction = lass_reg_model.predict(X_train)
```

```
[98]: error_score = metrics.r2_score(y_train, training_data_prediction)
print("R squared Error :", error_score)
```

R squared Error : 0.8424480718240743


```
[99]: plt.scatter(y_train, training_data_prediction)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual Prices vs Predicted Prices")
plt.show()
```



```
[100]: #prediction on training data
test_data_prediction = lass_reg_model.predict(X_test)
```

```
[102]: # R squared Error
error_score = metrics.r2_score(y_test, test_data_prediction)
print("R squared error : ", error_score)
```

R squared error : 0.8709763132343395

```
[104]: plt.scatter(y_test, test_data_prediction)
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual Prices vs Predicted Prices")
plt.show()
```



[]: